

```
In [1]: #This jupyter notebook is prepared by Jason Saini
```

```
In [2]: # import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import sklearn
from scipy import stats
import missingno as msno
```

```
In [3]: #import the data to a dataframe and show how many rows and columns it has
df = pd.read_csv('hrdata.csv')
print("Number of rows = " + str(df.shape[0]))
print("Number of columns = " + str(df.shape[1]))
```

Number of rows = 21287
Number of columns = 18

```
In [4]: #call the describe method of dataframe to see some summary statistics of the numerical columns
df.describe()
```

Out[4]:

	Unnamed: 0	rec_num	enrollee_id	city_development_index	training_hours	target	city_development_matrices
count	21287.000000	21287.000000	21287.000000	21287.000000	21287.000000	19158.000000	21287.000000
mean	10643.000000	10644.000000	16873.983652	0.828462	65.328510	0.249348	8.284615
std	6145.171926	6145.171926	9612.131237	0.123537	60.075201	0.432647	1.235365
min	0.000000	1.000000	1.000000	0.448000	1.000000	0.000000	4.480000
25%	5321.500000	5322.500000	8554.500000	0.739000	23.000000	0.000000	7.390000
50%	10643.000000	10644.000000	16967.000000	0.903000	47.000000	0.000000	9.030000
75%	15964.500000	15965.500000	25161.500000	0.920000	88.000000	0.000000	9.200000
max	21286.000000	21287.000000	33380.000000	0.949000	336.000000	1.000000	9.490000

```
In [5]: df.dtypes
```

Out[5]:

```
Unnamed: 0          int64
rec_num           int64
enrollee_id       int64
city              object
city_development_index  float64
gender            object
relevent_experience object
enrolled_university object
education_level   object
major_discipline  object
experience         object
company_size      object
company_type      object
last_new_job      object
training_hours    int64
target           float64
state            object
city_development_matrices float64
dtype: object
```

Interesting and good to know statistics:

- All counts besides target (19158) are exactly the same (21287)
- Max value == count for rec_num
- city_development_matrices == city_development_index * 10

```
In [6]: # show top 5 rows of dataframe
df.head()
```

Out[6]:

	Unnamed: 0	rec_num	enrollee_id	city	city_development_index	gender	relevent_experience	enrolled_university	education_level	major_discipline	experience	company_size
0	0	1	8949	city_103	0.920	Male	Has relevent experience	no_enrollment	Graduate	STEM	>20	None
1	1	2	29725	city_40	0.776	Male	No relevent experience	no_enrollment	Graduate	STEM	15	50
2	2	3	11561	city_21	0.624	NaN	No relevent experience	Full time course	Graduate	STEM	5	None
3	3	4	33241	city_115	0.789	NaN	No relevent experience	NaN	Graduate	Business Degree	<1	None

Unnamed: 0 rec_num enrollee_id city city_development_index gender relevent_experience enrolled_university education_level major_discipline experience company_s

```
In [7]: # Show last 5 rows of dataframe
df.tail()
```

Out[7]:

	Unnamed: 0	rec_num	enrollee_id	city	city_development_index	gender	relevent_experience	enrolled_university	education_level	major_discipline	experience	comp	
	21282	21282	21283	1289	city_103		0.920	Male	No relevent experience	no_enrollment	Graduate	Humanities	16
	21283	21283	21284	195	city_136		0.897	Male	Has relevent experience	no_enrollment	Masters	STEM	18
	21284	21284	21285	31762	city_100		0.887	Male	No relevent experience	no_enrollment	Primary School	NaN	3
	21285	21285	21286	7873	city_102		0.804	Male	Has relevent experience	Full time course	High School	NaN	7
	21286	21286	21287	12215	city_102		0.804	Male	Has relevent experience	no_enrollment	Masters	STEM	15

```
In [8]: # list all numerical columns
numerical_features = df.select_dtypes(include = np.number)
numerical_features.columns
```

Out[8]: Index(['Unnamed: 0', 'rec_num', 'enrollee_id', 'city_development_index', 'training_hours', 'target', 'city_development_matrices'], dtype='object')

```
In [9]: #list all categorical columns
categorical_features = df.select_dtypes(exclude = np.number)
categorical_features.columns
```

Out[9]: Index(['city', 'gender', 'relevent_experience', 'enrolled_university', 'education_level', 'major_discipline', 'experience', 'company_size', 'company_type', 'last_new_job', 'state'], dtype='object')

Examine Missing Values

```
In [10]: #Show a list with column wise count of missing values and display the list in count wise descending order
df.isnull().sum().sort_values(ascending = False)
```

Out[10]:


company_type	6774
company_size	6560
gender	5016
major_discipline	3125
target	2129
education_level	512
last_new_job	463
enrolled_university	417
experience	70
state	0
training_hours	0
Unnamed: 0	0
rec_num	0
relevent_experience	0
city_development_index	0
city	0
enrollee_id	0
city_development_matrices	0
dtype: int64	

```
In [11]: # Show a list with column wise percentage of missing values and display the list in percentage wise descending order'
df.isna().sum().sort_values(ascending = False) * 100 / len(df)
```

```
Out[11]: company_type          31.822239
company_size          30.816931
gender                23.563677
major_discipline      14.680321
target               10.001409
education_level        2.405224
last_new_job          2.175036
enrolled_university    1.958942
experience             0.328839
state                 0.000000
training_hours         0.000000
Unnamed: 0            0.000000
rec_num              0.000000
relevent_experience    0.000000
city_development_index 0.000000
city                 0.000000
enrollee_id           0.000000
city_development_matrices 0.000000
dtype: float64
```

```
In [12]: # Display a bar plot to visualize only the columns with missing values and their count
null_df = df.isnull().sum()
null_df = null_df[null_df > 0]
null_df.sort_values(inplace = True)
plt.xlabel("Column")
plt.ylabel("Number of missing values")
null_df.plot.bar()
```

```
Out[12]: <AxesSubplot:xlabel='Column', ylabel='Number of missing values'>
```




```
In [13]: #Use missingno's bar plot
msno.bar(df, color = 'blue', sort = "ascending" )
```

```
Out[13]: <AxesSubplot:>
```


```
#Use missingno's matrix plot w/ 200 sample
msno.matrix(df)
```

<AxesSubplot:>



```
# Use missingno's heatmap
msno.heatmap(df, cmap = "Reds")
```

<AxesSubplot:>







Interpret any interesting information you found in the heatmap and any one plot:

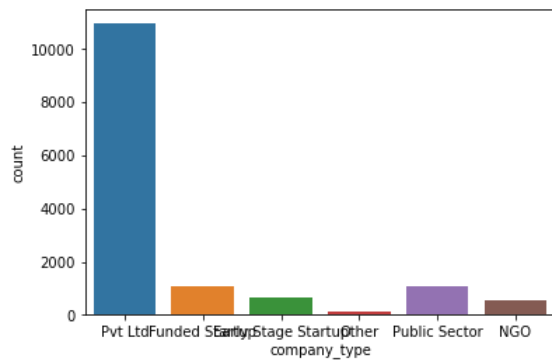




- matrix plot shows us that company_size & company_type are the two columns with highest missing values
- heatmap shows us that there is a strong positive correlation between company size and company type (might need to find a way to substitute missing values for higher accuracy before engineering a new feature)


Understanding Categorical attributes

```
In [16]: categorical_columns = df.select_dtypes(exclude = np.number)
```

```
In [18]: #Use seaborn bar plot for the categorical feature to see different values and count
for cat in categorical_columns:
    sns.countplot(x = cat, data = df)
    plt.show()
```






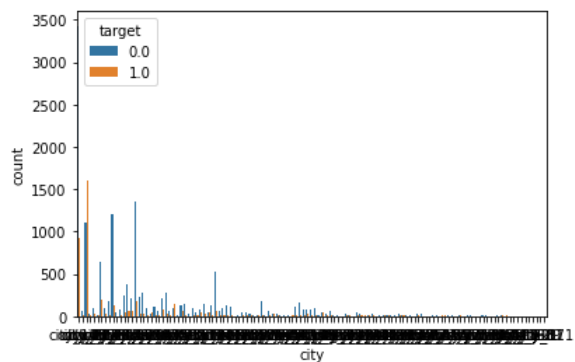










Note: countplot is a type of barplot that shows counts of each category, and its implementation is much cleaner

In [17]:

```
#Use seaborn countplot for the categorical feature against the values of the target
for cat in categorical_columns:
    sns.countplot(x = cat, hue = df["target"], data = df)
    plt.show()
```







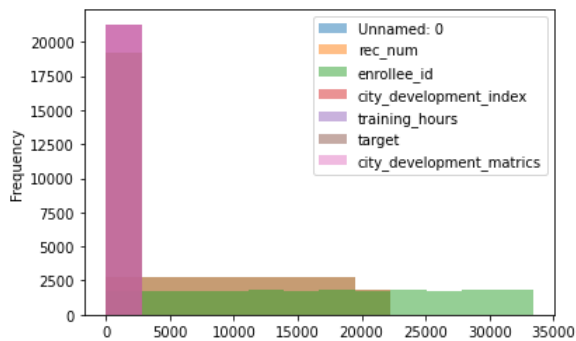
Interpret any interesting information and any information that might help you to make any decision on combining, removing, or adding features based on that, or any resampling maybe needed.

- We clearly need to resample for states as there is only one (CA)
- A lot of data is skewed to the left so we there might be a bias
- Might need to resample majors as there is a clear bias towards STEM

Understanding Numerical Attributes

```
In [19]: #Plot their distributions using histogram
numerical_features.plot.hist(bins = 12,alpha = .5)
```

```
Out[19]: <AxesSubplot:ylabel='Frequency'>
```




```
In [20]: #Plot the distribution using seaborn distplot
sns.distplot(numerical_features, hist = False)
```

C:\Users\Admin\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:ylabel='Density'>
```



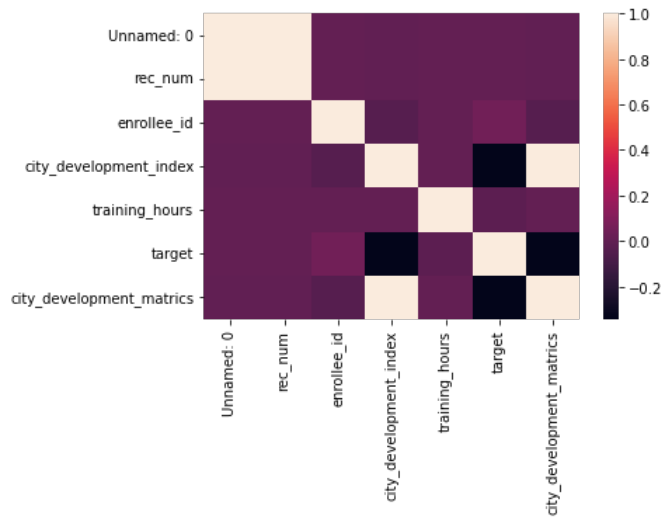
Interpret any interesting information:

- training_hours and city_development_index have much higher frequencies than the rest of the features

Correlation

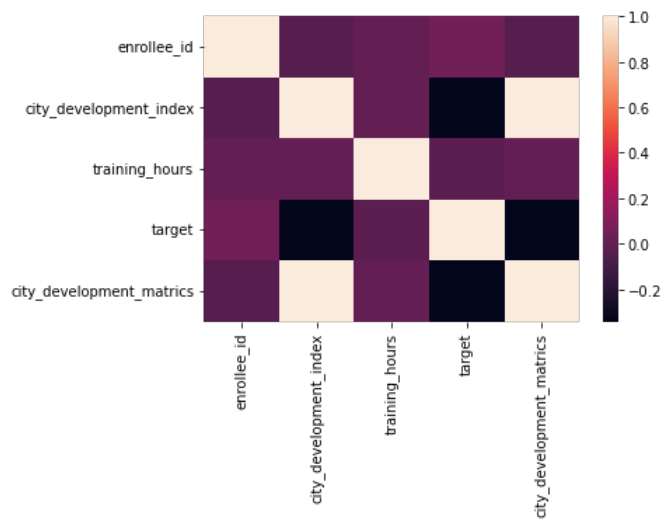
```
In [21]: #For the numerical attributes, use heatmap to show the correlation
sns.heatmap(numerical_features.corr())
```

```
Out[21]: <AxesSubplot:>
```








```
In [30]: # Heatmap with shortlisted columns
sns.heatmap(numerical_features.drop(labels = ["Unnamed: 0", "rec_num"], axis = 1).corr())
```



Out[30]: <AxesSubplot:>



```
In [23]: #Show scatter plots between columns to show the relationships with the target
for numerical_col in numerical_features:
    sns.relplot(data = numerical_features, y = numerical_col, x = "target", hue = numerical_col)
    plt.show()
```










Interpret and explain any finding and next course of action:





- We should re-evaluate the range for our target it is hard to distinguish when there is only 0 and 1
- We can remove rec_num feature as it has no correlation with our target value

Outliers

In [24]:


```
#Use boxplot or any other strategies to find outliers
for numerical_col in numerical_features:
    sns.boxplot(data = numerical_features, x = numerical_col)
    plt.show()
```





In [26]: `categorical_features["experience"].value_counts().plot(kind='bar')`

Out[26]: <AxesSubplot:>



What are the different values of experience, can you categorize them in to 0, 1, and 2?

- There are too many different values of experience to categorize in to 0,1 and 2 unless we utilize data ranges.

Summary and discussion: Finally after all the above EDA, summarize your finding, next course of action such as we may need to transform distribution because of right skew etc, need to remove a particular columns for any reasons, remove records for any reasons, need to rebalance data and what are the rebalancing options (if needed),

and any other finding.

- One part of our feature engineering should be to combine city_development_metrics and city_development_index (or we can remove one of them for simplicity)
- We should resample company_size & company_type as there are a lot of missing values
- Company_size & company_type can possibly be engineered into a single feature(or one can be dropped) but we cannot be sure until we resample the missing values (or fill them in with mean,median, etc).
- We need to resample states because our data is biased to CA (we only have data from one state)
- We might also have to resample our data for majors as there is a clear skew to STEM majors which may result in bias
- We should re-evaluate the range for our target it is hard to distinguish when there is only 0 and 1
- We can remove rec_num feature as it has no correlation with our target value